

JEN SEMLER

jen.semmler@philosophy.ox.ac.uk | www.jensemmler.com

RESEARCH AREAS

SPECIALIZATION	Philosophy of AI, practical ethics, normative ethics
COMPETENCE	Philosophy of mind, philosophy of action, experimental philosophy
AFFILIATIONS	Oxford Institute for Ethics in AI Oxford BioXPhi Lab ANU Machine Intelligence and Normative Theory Lab

EDUCATION

DPhil in Philosophy University of Oxford	2025 (expected)
Dissertation: <i>Artificial Moral Agency</i>	
Supervisors: Carissa Véliz and Alison Hills	
MPhil in Philosophy University of Cambridge	2021
Thesis: <i>The Right to Create Future Generations</i>	
MA in Medieval Icelandic Studies University of Iceland	2020
Thesis: <i>Is Njáll a Compatibilist? Aristotelian Voluntary Action in Njáls Saga</i>	
BA in Philosophy and Public Policy Duke University	2019
Magna cum laude; economics minor	

PUBLICATIONS

JOURNAL ARTICLES [†]	Recent Experimental Work on ‘Ought’ Implies ‘Can’	2019
	(with P. Henne)	
	<i>Philosophy Compass</i>	
	Against Some Recent Arguments for ‘Ought’ Implies ‘Can’	2019
	(with P. Henne, V. Chituc, F. de Brigard, and W. Sinnott-Armstrong)	
	<i>Philosophia</i>	
TEXTBOOK CONTRIBUTIONS	Corporate Human Rights Obligations	Under Contract
	case study for the “Human Rights” chapter in	
	<i>Issues in Political Theory</i> (edited by R. Jubb and P. Tomlin)	
	Oxford University Press	

[†] I share first authorship with P. Henne on both co-authored papers.

Lockdowns

Under Contract

case study for the “Liberty” chapter in
Issues in Political Theory (edited by R. Jubb and P. Tomlin)
 Oxford University Press

AWARDS

Oxford-Reuben Graduate Scholarship (full fees and living expenses for the DPhil)	2021-25
Global Priorities Fellowship (£5,000) <i>Forethought Foundation</i> in conjunction with the Early Career Conference Programme at the Global Priorities Institute (Oxford)	2022
Kate Bertram Prize (£100) <i>Lucy Cavendish College, Cambridge</i> for passing with distinction in the MPhil	2021
American Friends of Cambridge University Studentship (£2,500) <i>Lucy Cavendish College, University of Cambridge</i>	2020
Phi Beta Kappa Society <i>Duke University</i>	2019

TALKS (* indicates invited)

Moral Agents Unlike Us	
– Workshop on Partiality, Relationships, and AI <i>Ludwig Maximilian University of Munich (LMU)</i>	2024
Two Types of Moral Agency	
– Open Minds XVII <i>University of Manchester</i>	2024
Minimal Artificial Moral Agency	
– Agency and Intentions in AI Conference <i>University of Göttingen</i>	2024
Artificial Moral Behavior	
– Ethics and Computing Colloquium* <i>Cornell University</i>	2024
Locating Consciousness in Moral Agency	
– Machine Intelligence and Normative Theory Lab <i>Australian National University</i>	2023
AI and Moral Agency Attributions	
– European Experimental Philosophy Conference <i>University of Zurich</i>	2023
– Moral Psychology of AI Conference <i>University of Kent</i>	2023
– Oxford BioXPhi Lab <i>University of Oxford</i>	2023

Types of Artificial Moral Agency

- Philosophy, AI, and Society Doctoral Colloquium | *University of Oxford* 2023

Can AI be a Genuine Source of Moral Action?

- Oxford AI Society Mini-Conference* | *University of Oxford* 2024

Artificial Non-sentient Moral Agency

- AAI/ACM Conference on Artificial Intelligence, Ethics, and Society (AIES) Student Track Program | *University of Oxford* 2022

People Infer Temporal Order and Causal Structure from Causal Judgments

- European Experimental Philosophy Conference | *University of Granada* 2022
- Society for Philosophy and Psychology & European Society for Philosophy and Psychology (SPP & ESPP) Joint Conference | *University of Milan* 2022

Against the Epistemic Response to Moral Luck

- Free Will, Moral Responsibility, and Agency Graduate Conference | *Florida State University* 2021

TEACHING

PRIMARY	Practical Ethics	2024
INSTRUCTOR [‡]	Moral Agency in Nonhumans (with Virginie Simoneau-Gilbert)	2024
	Ethics of AI and Digital Technology	2022-23
	Ethical Computing in Practice	2023
	AI Ethics and Public Health	2023
	AI Ethics and Warfare	2023
	Ethics	2023
SEMINARS	Why Be Good? An Introduction to Ethics	2023-24
	<i>Exeter College Summer Programme, University of Oxford</i>	
TEACHING ASSISTANT	Why Be Good? An Introduction to Ethics	2022-23
	<i>Michael Hannon, Exeter College Summer Programme, University of Oxford</i>	

[‡] These courses were taught as Oxford-style tutorials: weekly, hour-long teaching sessions consisting of a tutor and 1-2 students. Tutors set a syllabus with readings and essay questions for each session. One exception: “Moral Agency in Nonhumans” was taught as a lecture series.

	Ethics (upper-level); Moral Philosophy (lower-level) <i>Jean Baccelli, University of Oxford</i>	2022-23
	Ethical Computing in Practice <i>Milo Phillips-Brown, University of Oxford</i>	2022
OTHER TEACHING	Moral Philosophy Seminar Facilitator <i>Weidenfeld-Hoffman Trust Leadership Programme, University of Oxford</i>	2023-24
	Oxford PPE UNIQ Summer School Teacher <i>University of Oxford</i>	2023
	Paradoxes: An Exploration of Problems in Philosophy and Beyond <i>Duke University</i>	2018-2019
	Volunteer high school tutor <i>Durham, NC</i>	2015-19
	Freelance philosophy tutor <i>virtual</i>	2017
	K-12 academic tutor <i>Turning Point Tutoring, NY</i>	2013-16

BROADER ENGAGEMENT

PRESENTATIONS	Invited Talks and Panels	
	– Artificial Moral Agency <i>University of Oxford Department of Continuing Education, “Thinking Machines and the Ethics of AI”</i>	2025
	– Can AI Be Moral? <i>Magdalen College School, Oxford</i>	2024
	– AI Safety and Diversity, Equity, and Inclusion: A FemTech Perspective <i>Oxford FemTech Society</i>	2024
	– Making Moral Machines <i>Oxford University Computer Society and Hertford Politics and Economics Society</i>	2022
	Other Talks, Panels, and Exhibitions	
	– What Even Is Philosophy Research? <i>Reuben College</i>	2024
	– The Moral Mind <i>Oxford Museum of Natural History Brain Awareness Week, Art and Science Exhibition</i>	2023
	– What Does the Future Hold for Humanity? <i>Reuben College</i>	2022
	– AI and the Moral Community <i>Reuben College</i>	2022
	– Experimental Philosophy <i>Reuben College</i>	2022
	– Jen Semler on AI, Icelandic Studies, Women in Philosophy, and Publishing <i>The Philosopher’s Nest</i> podcast listen here	2022

RELEVANT EXPERIENCE

RESEARCH ASSISTANTSHIPS	Milo Phillips-Brown (ethics of technology) <i>University of Oxford, University of Edinburgh</i>	2022-24
	Project Vox <i>Duke University</i>	2018-19

	Moral Attitudes and Decision-Making Lab <i>Duke University</i>	2016-19
EMPLOYMENT	Philosophy Undergraduate Admissions Interviewer <i>Jesus College, University of Oxford</i>	2021-23
	Research Programs Intern <i>National Endowment for the Humanities, Washington, D.C.</i>	2018
SERVICE		
CONFERENCES	Organizing Committee Member	
	– Oxford Women and Gender Minorities in Philosophy Internal Workshop	2022
	– Oxford Graduate Philosophy Conference	2022
	– Háskóli Íslands Student Conference on the Medieval North (cancelled due to COVID-19)	2020
	Session Chair and/or Assistant	
	– Oxford Graduate Philosophy Conference	2023
	– Responsibility and Autonomy in Artificial Intelligence Conference (Oxford)	2022
	– Cambridge Graduate Conference on the Philosophy of Mathematics and Logic	2021
REPRESENTATION	Women and Gender Minorities Graduate Philosophy Student Representative <i>University of Oxford</i>	2023-24
	Graduate Director for Ethics and Values Theme <i>Reuben College</i>	2022-24
REVIEWER	Journals: <i>Philosophical Psychology; Philosophy & Technology</i>	
	Conferences: <i>Oxford Graduate Philosophy Conference; Fairness, Accountability, and Transparency Conference (ACM FAccT)</i>	
MISCELLANEOUS	Consent Workshop Facilitator <i>Reuben College</i>	2023-24
	Judge <i>Ethics Cup (formerly JS Mill Cup), held online due to COVID-19</i>	2021

REFERENCES

Carissa Véliz

Associate Professor of Philosophy
University of Oxford
carissa.veliz@philosophy.ox.ac.uk

Alison Hills

Professor of Philosophy
University of Oxford
alison.hills@sjc.ox.ac.uk

Milo Phillips-Brown

Lecturer in Philosophy
University of Edinburgh
milopb@ed.ac.uk

Joanna Demaree-Cotton

Research Fellow in Moral Psychology
University of Oxford
joanna.demaree-cotton@philosophy.ox.ac.uk

Michael Hannon (teaching)

Associate Professor of Philosophy
University of Nottingham
michael.hannon@nottingham.ac.uk

DISSERTATION ABSTRACT: ARTIFICIAL MORAL AGENCY

Suppose Aasha destroys Dottie's house. We might have some questions—whether, for instance, Aasha committed a moral wrong or whether Aasha is morally responsible. The answers to these questions depend, in part, on what kind of entity Aasha is. Our evaluation of this scenario differs if Aasha is a strong gust of wind, or a giraffe, or a human. Specifically, our assessment turns on whether Aasha is a *moral agent*. My dissertation considers what follows if Aasha is an AI system.

On my novel account of moral agency, AI systems can be moral agents in principle, but existing AI systems fail to meet the necessary conditions. As such, our use of AI in moral decision-making should be limited. Moreover, even if future AI systems are genuine moral agents, artificial moral agents will be different from human moral agents in normatively significant ways.

Part I—Towards a Theory of Moral Agency—develops a theoretical account of moral agency. In Chapter 1, “Locating Consciousness in Moral Agency,” I preempt a common objection to the prospect of artificial moral agency, namely that AI systems cannot be moral agents because they lack consciousness. I argue that phenomenal consciousness is not necessary for moral agency. In Chapter 2, “Two Types of Moral Agency,” I propose and defend a distinction: *deontic moral agents* are appropriate subjects of deontic evaluations—their actions can be described as morally wrong—and *responsible moral agents* are appropriate subjects of responsibility ascriptions—they are fully-fledged moral agents. This distinction illuminates difficult cases of moral agency (e.g., children, psychopaths, and nonhuman animals) as well as instances of genuine responsibility gaps (cases in which a moral wrong has been committed for which no one is morally responsible).

Part II—Prospects of Artificial Moral Agency—evaluates the extent to which AI systems are moral agents. I take a technically grounded approach: I consider whether existing machine learning methods and empirical results support classifying AI systems as moral agents. In Chapter 3, “Artificial ‘Agents’ are Not Agents,” I argue that AI systems lack the kind of agency required for moral agency—namely, the capacity for intentional action. Against recent claims to the contrary, I argue that AI systems, especially large language models and reinforcement learning systems, lack mental states. In Chapter 4, “Artificial ‘Agents’ are Not Moral,” I argue that even if existing AI systems were agents, they are far from instantiating the relevant capacities required for deontic and responsible moral agency (which I outline in Chapter 2). AI systems lack responsiveness to moral reasons and moral understanding.

Part III—Using Artificial (non) Moral Agents—considers how the moral agency, or lack thereof, of AI systems bears on how we use those systems in moral decision-making. In Chapter 5, “Artificial Moral Behavior,” I argue that delegating moral decisions to AI systems is wrong—even if the AI system's outputs are reliable and accurate—because doing so replaces moral *actions* with, at best, moral *behaviors*. We have both intrinsic and instrumental reasons to avoid shrinking the moral domain in this way. In Chapter 6, “Moral Agents Unlike Us,” I argue that even if AI systems qualify for responsible moral agency, they are different from human moral agents in morally significant ways. Even though a lack of consciousness is no barrier to moral agency (as I argue in Chapter 1), a lack of consciousness *is* a barrier to playing certain roles in the moral community. Moral agency isn't all that matters.