

JEN SEMLER

jensemlephilosophy@gmail.com | www.jensemle.com

RESEARCH AREAS

SPECIALIZATION	Philosophy of AI, practical ethics, normative ethics
COMPETENCE	Philosophy of mind, philosophy of action, experimental philosophy
AFFILIATIONS	Oxford Institute for Ethics in AI Oxford Philosophical Moral Psychology Lab ANU Machine Intelligence and Normative Theory Lab

EMPLOYMENT

Postdoctoral Associate | Cornell Tech (Digital Life Initiative) 2025—

EDUCATION

DPhil in Philosophy University of Oxford Dissertation: <i>On Artificial Moral Agency</i> Supervisors: Carissa Véliz and Alison Hills Examiners: Roger Crisp and David Shoemaker	2025
MPhil in Philosophy University of Cambridge Thesis: <i>The Right to Create Future Generations</i>	2021
MA in Medieval Icelandic Studies University of Iceland Thesis: <i>Is Njáll a Compatibilist? Aristotelian Voluntary Action in Njáls Saga</i>	2020
BA in Philosophy and Public Policy Duke University Magna cum laude; economics minor	2019

PUBLICATIONS

JOURNAL ARTICLES [†]	Moral Agency Without Consciousness Forthcoming in <i>Canadian Journal of Philosophy</i>
	A Timing Problem for Instrumental Convergence (with R. Southan and H. Ward) Forthcoming in <i>Philosophical Studies</i>

[†] Unless otherwise noted, all authors share first-authorship on co-authored papers.

Recent Experimental Work on ‘Ought’ Implies ‘Can’ 2019
(with P. Henne)
Philosophy Compass

Against Some Recent Arguments for ‘Ought’ Implies ‘Can’ 2019
(with P. Henne, V. Chituc, F. de Brigard, and W. Sinnott-Armstrong)
Philosophia
I share first-authorship with P. Henne

TEXTBOOK
CONTRIBUTIONS

Corporate Human Rights Obligations Under Contract
case study for the “Human Rights” chapter in
Issues in Political Theory (edited by R. Jubb and P. Tomlin)
Oxford University Press

Lockdowns Under Contract
case study for the “Liberty” chapter in
Issues in Political Theory (edited by R. Jubb and P. Tomlin)
Oxford University Press

AWARDS

Oxford-Reuben Graduate Scholarship (full fees and living expenses for the DPhil) 2021-25

Global Priorities Fellowship (£5,000) | *Forethought Foundation* 2022
in conjunction with the Early Career Conference Programme at the Global Priorities Institute (Oxford)

Kate Bertram Prize (£100) | *Lucy Cavendish College, Cambridge* 2021
for passing with distinction in the MPhil

American Friends of Cambridge University Studentship (£2,500) | 2020
Lucy Cavendish College, University of Cambridge

Phi Beta Kappa Society | *Duke University* 2019

TALKS (* indicates invited)

Moral Agents Unlike Us
– Workshop on Partiality, Relationships, and AI | *Ludwig Maximilian University of Munich (LMU)* 2024

Two Types of Moral Agency
– Open Minds XVII | *University of Manchester* 2024

Minimal Artificial Moral Agency

- Agency and Intentions in AI Conference | *University of Göttingen* 2024

Artificial Moral Behavior

- Ethics and Computing Colloquium* | *Cornell University* 2024
- Machine Intelligence and Normative Theory Lab | *Australian National University* 2024

Locating Consciousness in Moral Agency

- Machine Intelligence and Normative Theory Lab | *Australian National University* 2023

AI and Moral Agency Attributions

- European Experimental Philosophy Conference | *University of Zurich* 2023
- Moral Psychology of AI Conference | *University of Kent* 2023
- Oxford BioXPhi Lab | *University of Oxford* 2023

Types of Artificial Moral Agency

- Philosophy, AI, and Society Doctoral Colloquium | *University of Oxford* 2023

Can AI be a Genuine Source of Moral Action?

- Oxford AI Society Mini-Conference* | *University of Oxford* 2024

Artificial Non-sentient Moral Agency

- AAAI/ACM Conference on Artificial Intelligence, Ethics, and Society (AIES) Student Track Program | *University of Oxford* 2022

People Infer Temporal Order and Causal Structure from Causal Judgments

- European Experimental Philosophy Conference | *University of Granada* 2022
- Society for Philosophy and Psychology & European Society for Philosophy and Psychology (SPP & ESPP) Joint Conference | *University of Milan* 2022

Against the Epistemic Response to Moral Luck

- Free Will, Moral Responsibility, and Agency Graduate Conference | *Florida State University* 2021

TEACHING

PRIMARY	Practical Ethics	2024
---------	------------------	------

INSTRUCTOR [‡]	Moral Agency in Nonhumans (with Virginie Simoneau-Gilbert)	2024
	Ethics of AI and Digital Technology	2022-24
	Ethical Computing in Practice	2023
	AI Ethics and Public Health	2023
	AI Ethics and Warfare	2023
	Ethics	2023
THESIS SUPERVISION	Undergraduate thesis on AI and responsibility	2025
SEMINARS	Why Be Good? An Introduction to Ethics	2023-24
	<i>Exeter College Summer Programme, University of Oxford</i>	
TEACHING ASSISTANT	Why Be Good? An Introduction to Ethics	2022-23
	<i>Michael Hannon, Exeter College Summer Programme, University of Oxford</i>	
	Ethics (upper-level); Moral Philosophy (lower-level)	2022-23
	<i>Jean Baccelli, University of Oxford</i>	
	Ethical Computing in Practice	2022
	<i>Milo Phillips-Brown, University of Oxford</i>	
OTHER TEACHING	Teaching Ethics to Machines <i>Bitesize Ethics, Uehiro Oxford Institute</i>	2025
	Moral Philosophy Seminar Facilitator <i>Weidenfeld-Hoffman Trust Leadership Programme, University of Oxford</i>	2023-25
	Oxford PPE UNIQ Summer School Teacher <i>University of Oxford</i>	2023
	Paradoxes: An Exploration of Problems in Philosophy and Beyond <i>Duke University</i>	2018-2019
	Volunteer high school tutor <i>Durham, NC</i>	2015-19
	Freelance philosophy tutor <i>virtual</i>	2017
	K-12 academic tutor <i>Turning Point Tutoring, NY</i>	2013-16

BROADER ENGAGEMENT

PRESENTATIONS Invited Talks and Panels

[‡] These courses were taught as Oxford-style tutorials: weekly, hour-long teaching sessions consisting of a tutor and 1-2 students. Tutors set a syllabus with readings and essay questions for each session. One exception: “Moral Agency in Nonhumans” was taught as a lecture series.

- The Rise of Deepfake AI (panel) | *Reuben College* 2025
- Outsourcing to AI | *University of Oxford Department of Continuing Education, "Thinking Machines and the Ethics of AI"* 2025
- Can AI Be Moral? | *Magdalen College School, Oxford* 2024
- AI Safety and Diversity, Equity, and Inclusion: A FemTech Perspective | *Oxford FemTech Society* 2024
- Making Moral Machines | *Oxford University Computer Society and Hertford Politics and Economics Society* 2022

Other Talks, Panels, and Exhibitions

- Moral Responsibility and AI (talk) | *London Academy of Excellence, Tottenham* 2025
- What Even Is Philosophy Research? (talk) | *Reuben College* 2024
- The Moral Mind (exhibition) | *Oxford Museum of Natural History Brain Awareness Week* 2023
- What Does the Future Hold for Humanity? (panel) | *Reuben College* 2022
- AI and the Moral Community (talk) | *Reuben College* 2022
- Experimental Philosophy (talk) | *Reuben College* 2022
- Jen Semler on AI, Icelandic Studies, Women in Philosophy, and Publishing (interview) | *The Philosopher's Nest* podcast 2022

RELEVANT EXPERIENCE

RESEARCH ASSISTANTSHIPS	Milo Phillips-Brown (ethics of technology) <i>University of Oxford, University of Edinburgh</i>	2022-24
	Project Vox <i>Duke University</i>	2018-19
	Moral Attitudes and Decision-Making Lab <i>Duke University</i>	2016-19
EMPLOYMENT	Philosophy Undergraduate Admissions Interviewer <i>Jesus College, University of Oxford</i>	2021-23
	Research Programs Intern <i>National Endowment for the Humanities, Washington, D.C.</i>	2018

SERVICE

CONFERENCES	Organizing Committee Member	
	– Oxford Women and Gender Minorities in Philosophy Internal Workshop <i>University of Oxford</i>	2022
	– Oxford Graduate Philosophy Conference <i>University of Oxford</i>	2022
	– Háskóli Íslands Student Conference on the Medieval North (cancelled due to COVID-19) <i>University of Iceland</i>	2020

	Commenter	
	– Artificial Intelligence and Collective Agency Workshop <i>University of Oxford</i>	2025
	Session Chair and/or Assistant	
	– Oxford Graduate Philosophy Conference <i>University of Oxford</i>	2023
	– Responsibility and Autonomy in Artificial Intelligence Conference <i>University of Oxford</i>	2022
	– Cambridge Graduate Conference on the Philosophy of Mathematics and Logic <i>University of Cambridge</i>	2021
REPRESENTATION	Women and Gender Minorities Graduate Philosophy Student Representative <i>University of Oxford</i>	2023-24
	Graduate Director for Ethics and Values Theme <i>Reuben College</i>	2022-24
REVIEWER	Journals: <i>Philosophical Psychology</i> ; <i>Philosophy & Technology</i>	
	Conferences: <i>Oxford Graduate Philosophy Conference</i> ; <i>Fairness, Accountability, and Transparency Conference (ACM FAccT)</i> ; <i>Artificial Intelligence, Ethics, and Society Conference (AAAI/ACM AIES)</i>	
MISCELLANEOUS	Judge “ <i>Philosothon</i> ” at <i>Magdalen College School, Oxford</i>	2025
	Consent Workshop Facilitator <i>Reuben College</i>	2023-24
	Judge <i>Ethics Cup (formerly JS Mill Cup)</i> , held online	2021

REFERENCES

Carissa Véliz

Associate Professor of Philosophy | University of Oxford

carissa.veliz@philosophy.ox.ac.uk

Alison Hills

Professor of Philosophy | University of Oxford

alison.hills@sjc.ox.ac.uk

David Shoemaker

Wyn and William Y. Hutchinson Professor of Ethics and Public Life | Cornell University

dws267@cornell.edu

Roger Crisp

Professor of Moral Philosophy | University of Oxford

roger.crisp@uehiro.ox.ac.uk

Joanna Demaree-Cotton

Research Fellow in Moral Psychology | University of Oxford

joanna.demaree-cotton@uehiro.ox.ac.uk

Milo Phillips-Brown

Lecturer in Philosophy | University of Edinburgh

milopb@ed.ac.uk

Michael Hannon (teaching)

Associate Professor of Philosophy | University of Nottingham

michael.hannon@nottingham.ac.uk

DISSERTATION ABSTRACT: ON ARTIFICIAL MORAL AGENCY

Suppose Tommy destroys Marvin's house. We might have some questions—whether, for instance, Tommy committed a moral wrong or whether Tommy is morally responsible. The answers to these questions depend, in part, on what kind of entity Tommy is. Our evaluation of this scenario differs if Tommy is a strong gust of wind, or a giraffe, or a human. Specifically, our assessment turns on whether Tommy is a *moral agent*. My dissertation considers what follows if Tommy is an AI system. On my account of moral agency, AI systems can be moral agents in principle, but existing AI systems fail to meet the necessary conditions. As such, our use of AI in moral decision-making should be limited. Moreover, genuine artificial moral agents will be different from human moral agents in normatively significant ways.

Part I—Towards a Theory of Moral Agency—develops a theoretical account of moral agency. In Chapter 2, “Moral Agency Without Consciousness” I preempt an objection to the prospect of artificial moral agency, namely that AI systems cannot be moral agents because they lack consciousness. I argue that phenomenal consciousness is not necessary for moral agency. In Chapter 3, “Two Types of Moral Agency,” I propose and defend a distinction: *deontic moral agents* are appropriate subjects of deontic evaluations—their actions can be described as morally wrong—and *responsible moral agents* are appropriate subjects of responsibility ascriptions—they are fully-fledged moral agents. This distinction illuminates difficult cases of moral agency as well as instances of genuine responsibility gaps.

Part II—Prospects of Artificial Moral Agency—evaluates the extent to which AI systems are moral agents. I consider whether existing machine learning methods and empirical results support classifying existing AI systems, specifically large language models and reinforcement learning systems, as moral agents. In Chapter 4, “Artificial ‘Agents’ are Not Agents,” I argue that AI systems lack the kind of agency required for moral agency—namely, the capacity for intentional action—because they lack mental states on both interpretivist and representationalist views. In Chapter 5, “Artificial ‘Agents’ are Not Moral,” I argue that AI systems are far from instantiating the additional necessary capacities for deontic and responsible moral agency: AI systems lack responsiveness to moral reasons and moral understanding.

Part III—Using Artificial (non) Moral Agents—considers how the moral agency of AI systems, or lack thereof, bears on how we should use those systems in moral decision-making. In Chapter 6, “Artificial Moral Behavior,” I argue that delegating moral decisions to AI systems is wrong—even if the outputs are reliable and accurate—because doing so replaces moral *actions* with, at best, moral *behaviors*. In Chapter 7, “Moral Agents Unlike Us,” I argue that even if AI systems qualify for responsible moral agency, they are different from human moral agents in morally significant ways. While their lack of consciousness is no barrier to moral agency, it is a barrier to playing certain roles in the moral community. Moral agency is not all that matters.